

Capstone Project: Predict Future Sales

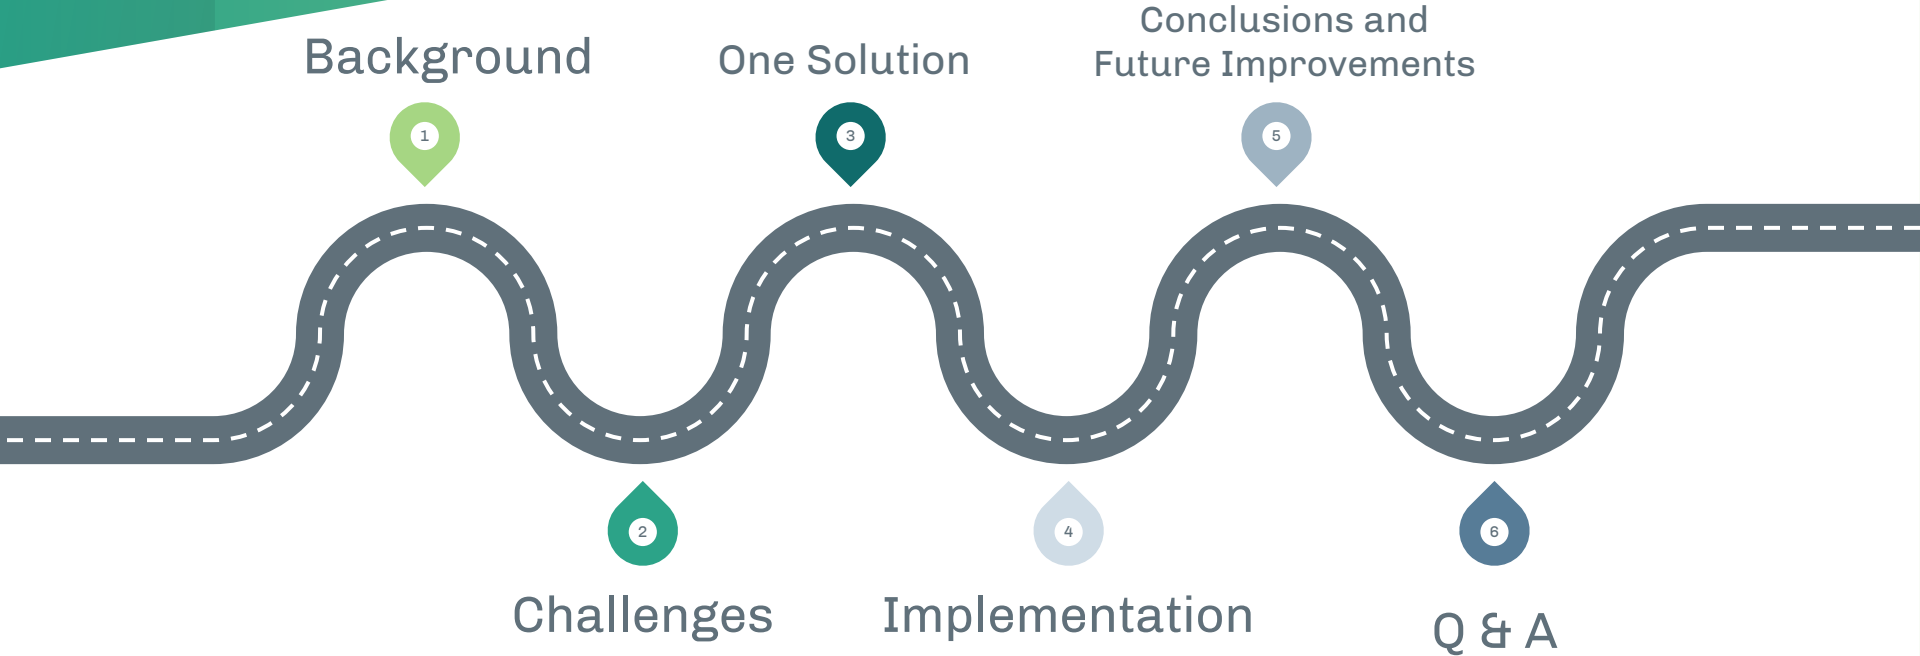
❖ Hierarchical Forecasting & Reconciliation

❖ Feng Yong He

Problem Statement

To build robust models to forecast monthly sales for each product in each retail store of 1C company, provided that most historical sales data are intermittent and sparse

Roadmap



Background



Predict Future Sales

Final project for "How to win a data science competition" Coursera course
Playground · 15598 Teams · 2 months to go

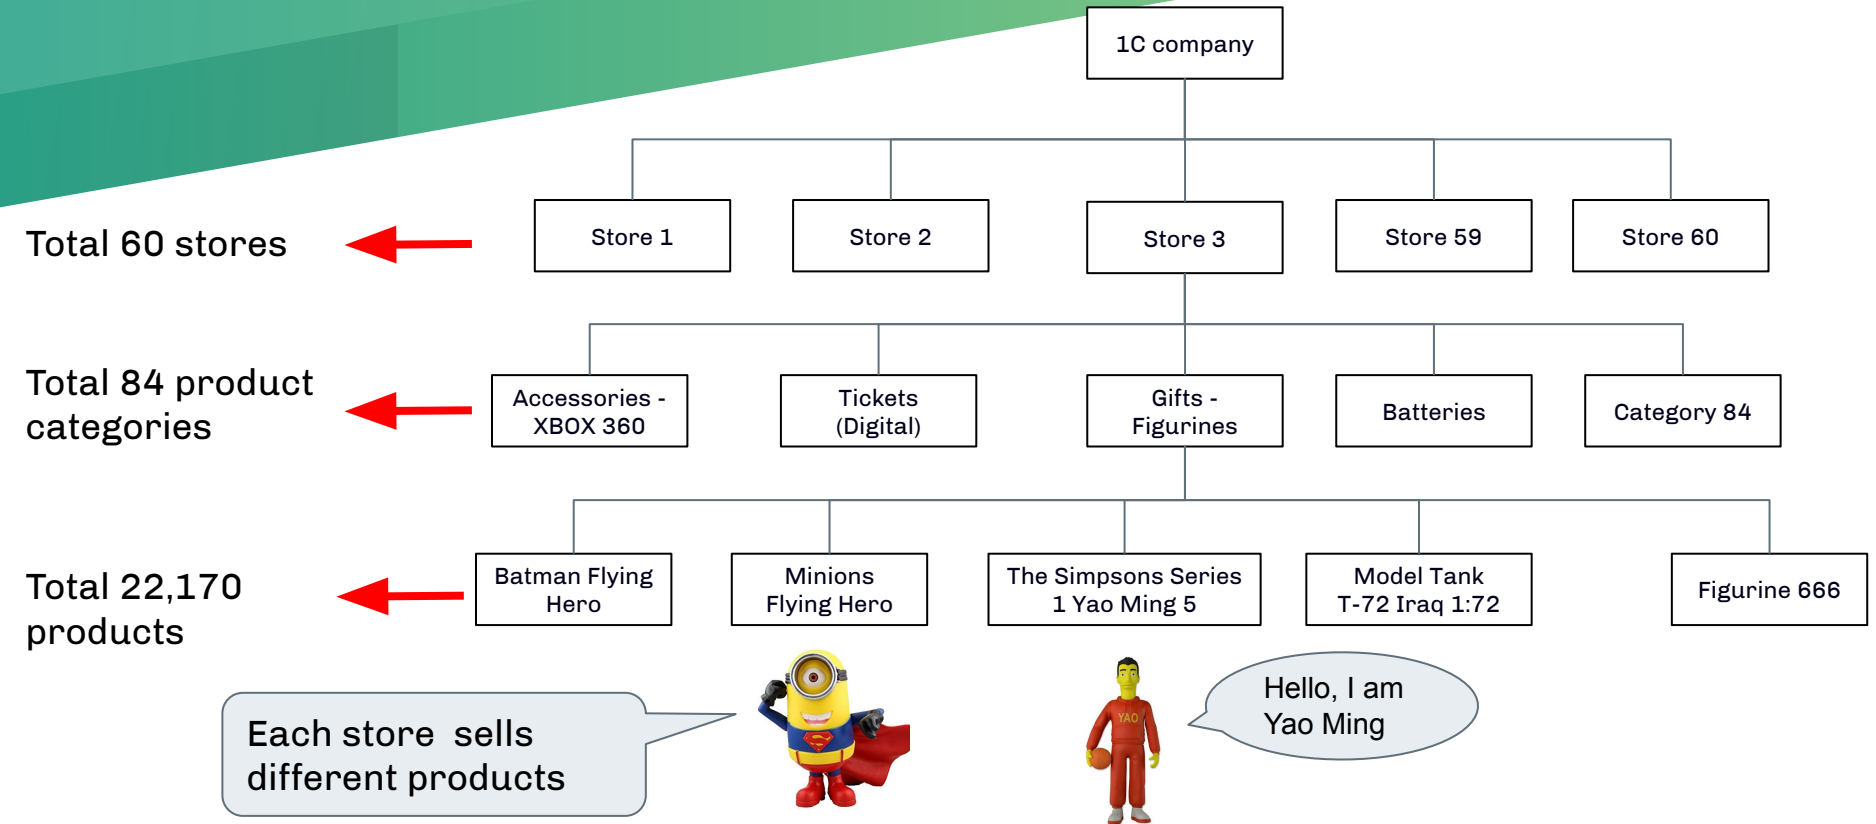
- One of Kaggle Active Competitions
- Dataset are provided by 1C Company

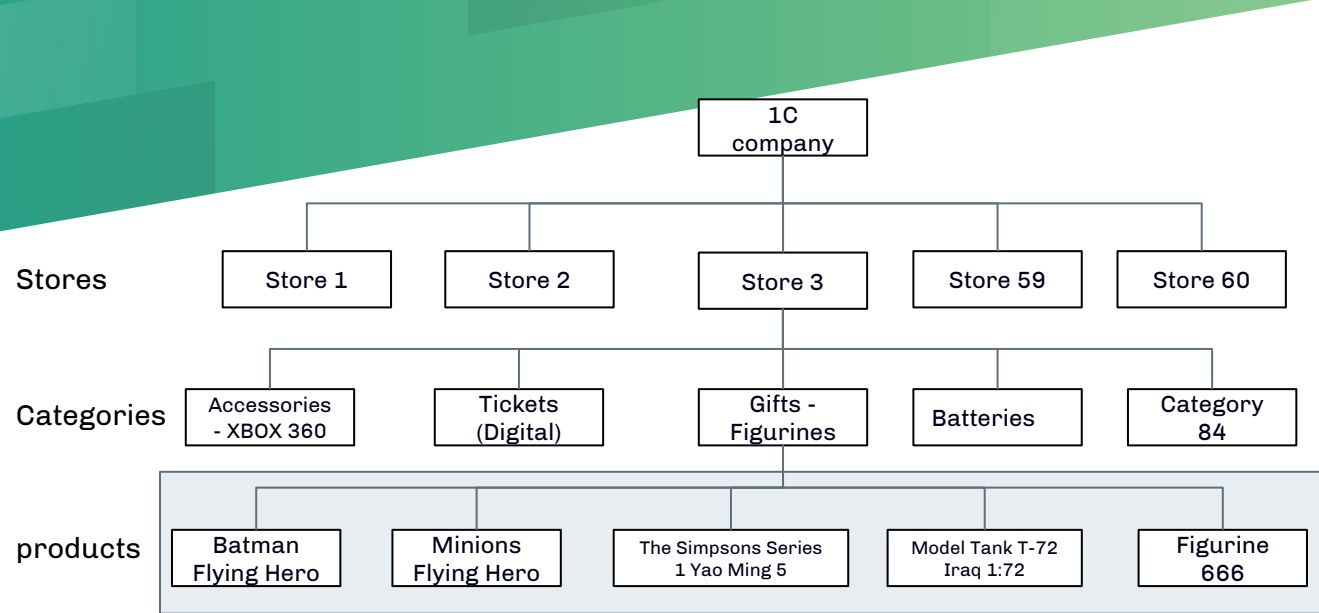
1C Company



- A Russian software developer, distributor and publisher based in Moscow.
- It is also running a retail chain selling computer software, related services, game consoles, video games, etc.

Retail Chain Structure





- ❑ We are given daily historical sales data at this level
- ❑ Our task is to predict monthly sales of each product at this level too

date	item_id	item_category_id	shop_id	item_price	item_cnt_day
2014-08-30	14263	38	37	199.00	1.00
2014-05-11	5459	55	25	349.00	1.00
2013-05-22	13102	4	15	999.00	1.00

Challenges I


- Massive datasets
 - 2.9 millions records
 - From Jan 2013 to Oct 2015
- Messy data
 - the list of stores and products slightly changed every month
 - Some stores have been shut down while some stores just opened for one month
 - Some stores have very similar name, and seems to be same store
 - Zhukovsky st. Chkalova 39m?
 - Zhukovsky st. Chkalova 39m²
 - Quite a lot new arrival products without any historical sales data

Challenges II

- Intermittent and sparse historical sales data
 - **Most products only have occasional transactions**
 - The sales time series data contains many embedded zeros
 - Has no clearly defined trend
 - Does NOT exhibit any seasonal behavior
 - Very difficult for a conventional time series model such as ARIMA to forecast.



I had 25 transactions in all shops in 2015!



I only had 1 transaction in 2015

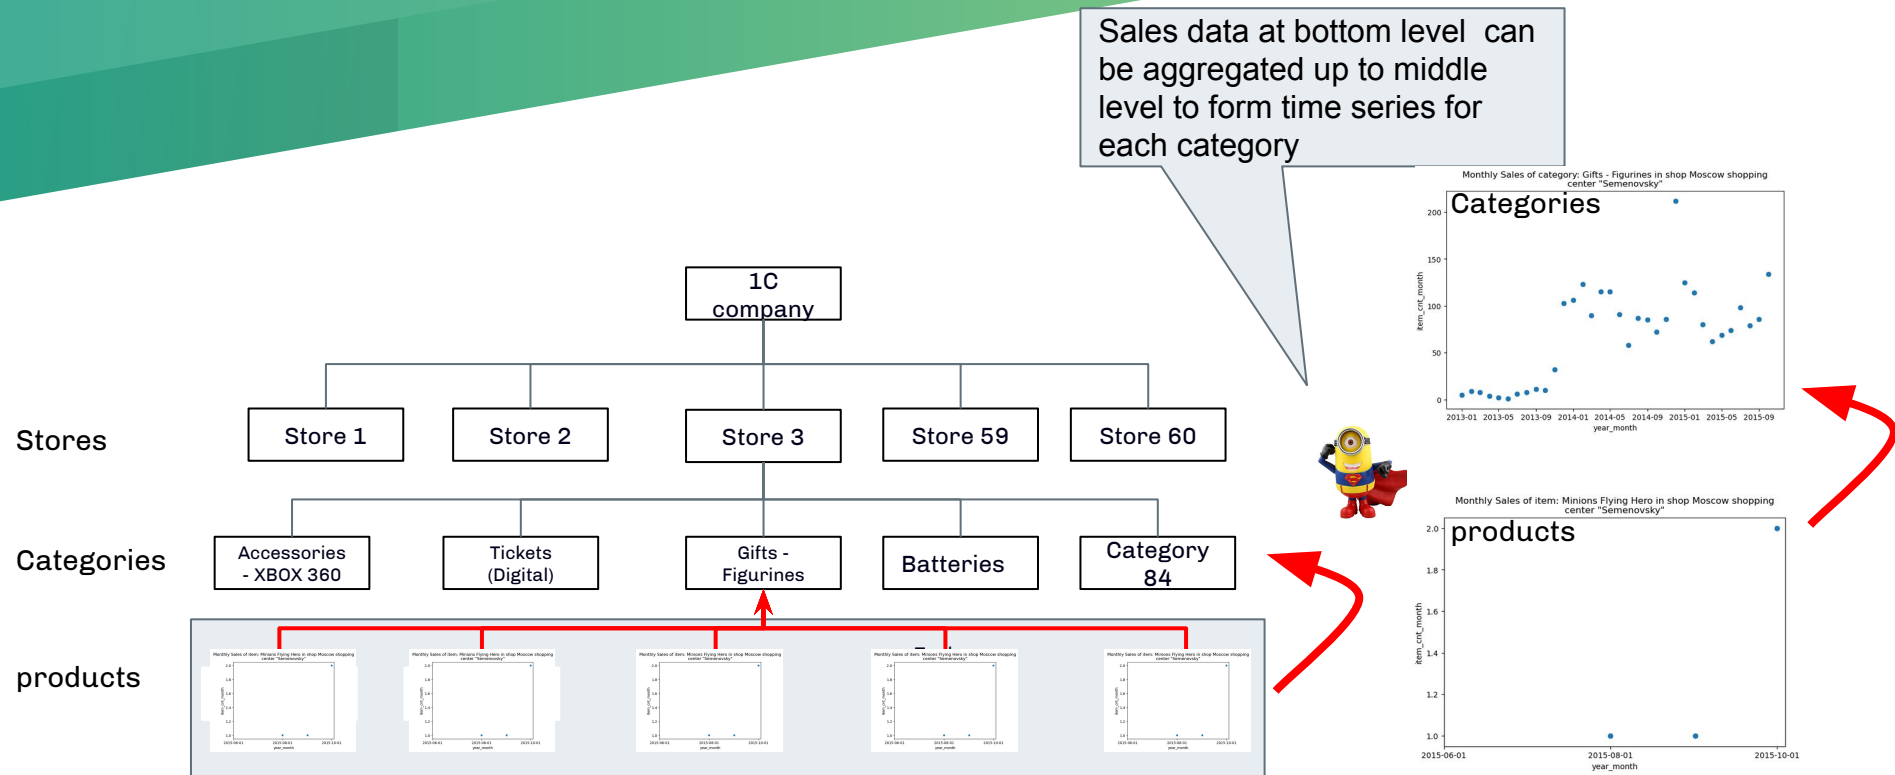


Hierarchical Time Series Forecasting & Reconciliation

-
- The diagram illustrates the three-step process of the proposed framework for hierarchical time series forecasting:
- Step 1: Data** - A hierarchical tree structure showing the relationship between different levels of aggregation. The root node is 'Total', which branches into 'A' and 'B'. 'A' further branches into 'AA' and 'AB', while 'B' branches into 'BA' and 'BB'.
 - Step 2: Machine learning models** - The data is processed by machine learning models (represented by gears) to generate forecasts. The output is a vector of forecasts: $\tilde{Y}_t = \begin{pmatrix} \tilde{Y}_{Total} \\ \tilde{Y}_A \\ \tilde{Y}_B \\ \tilde{Y}_{AA} \\ \tilde{Y}_{AB} \\ \tilde{Y}_{BA} \\ \tilde{Y}_{BB} \end{pmatrix}$.
 - Step 3: Reconciliation models** - The forecasts are processed by reconciliation models (represented by gears) to produce reconciled forecasts. The output is a vector of reconciled forecasts: $\hat{Y}_t = \begin{pmatrix} \hat{Y}_{Total} \\ \hat{Y}_A \\ \hat{Y}_B \\ \hat{Y}_{AA} \\ \hat{Y}_{AB} \\ \hat{Y}_{BA} \\ \hat{Y}_{BB} \end{pmatrix}$.



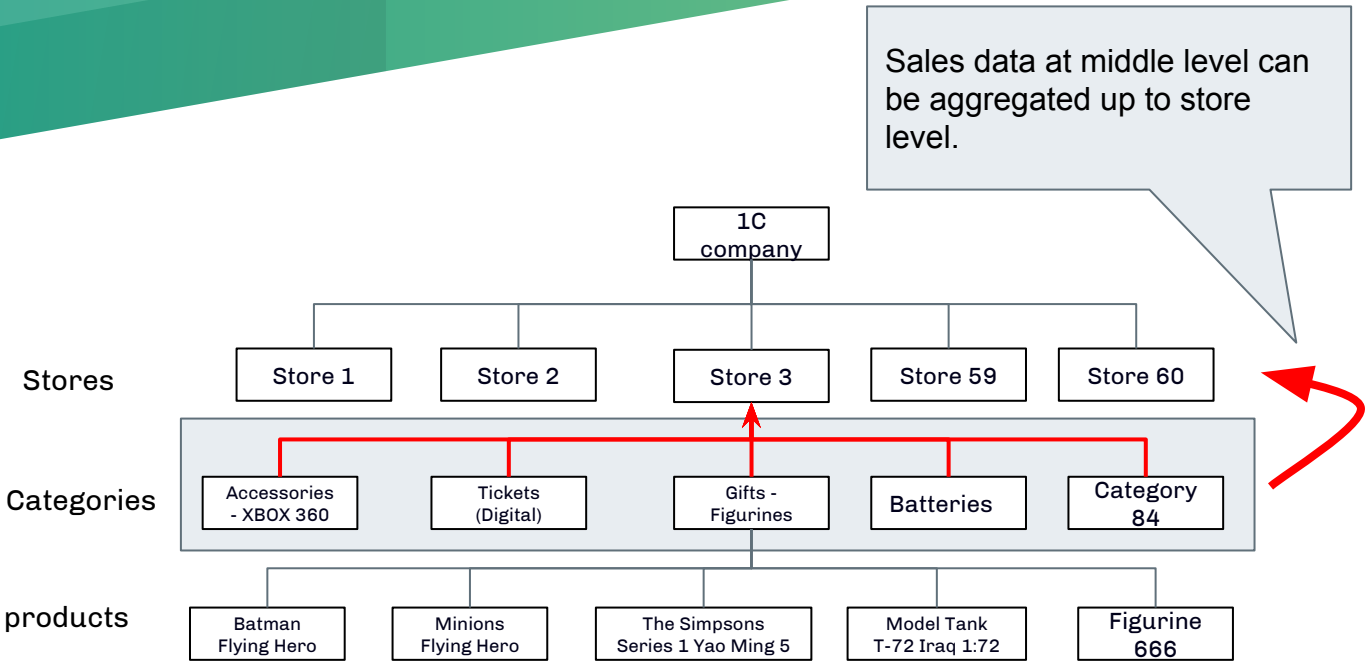
Step 1: Building Hierarchical Time Series:

Substep 1: Aggregation From Bottom Level



Step 1: Building Hierarchical Time Series:

Substep 2: Aggregation From Middle Level



Step 1: Building Hierarchical Time Series:



Done!

Hierarchical Time Series

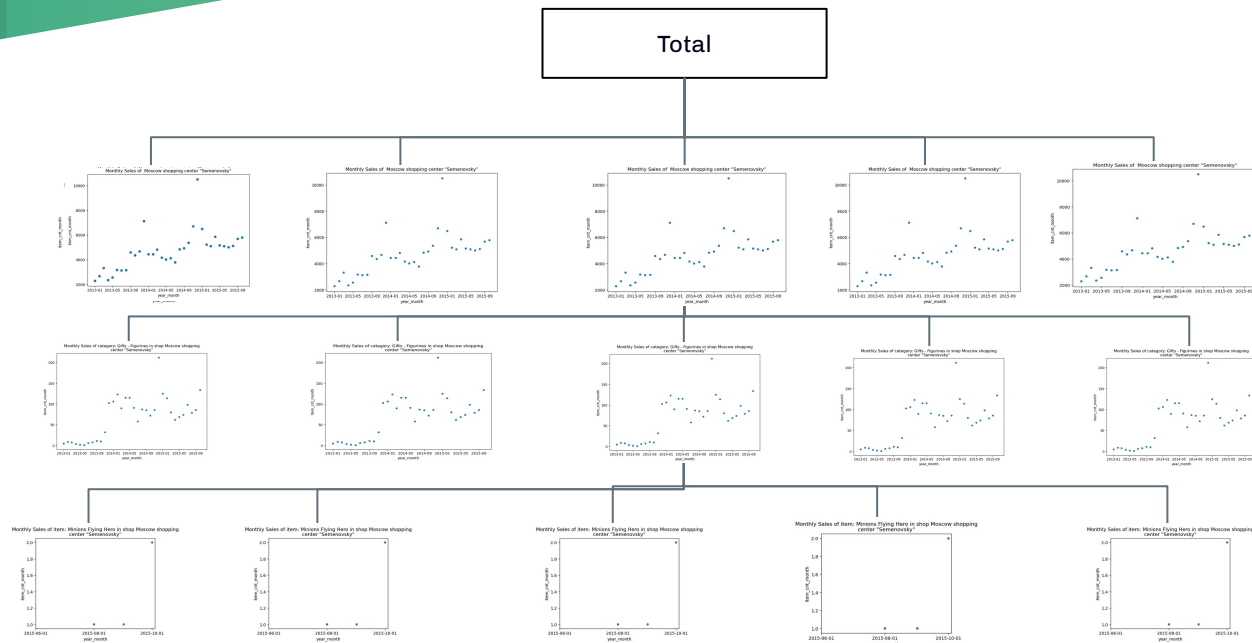
Stores

Categories

Products


Aggregate

12

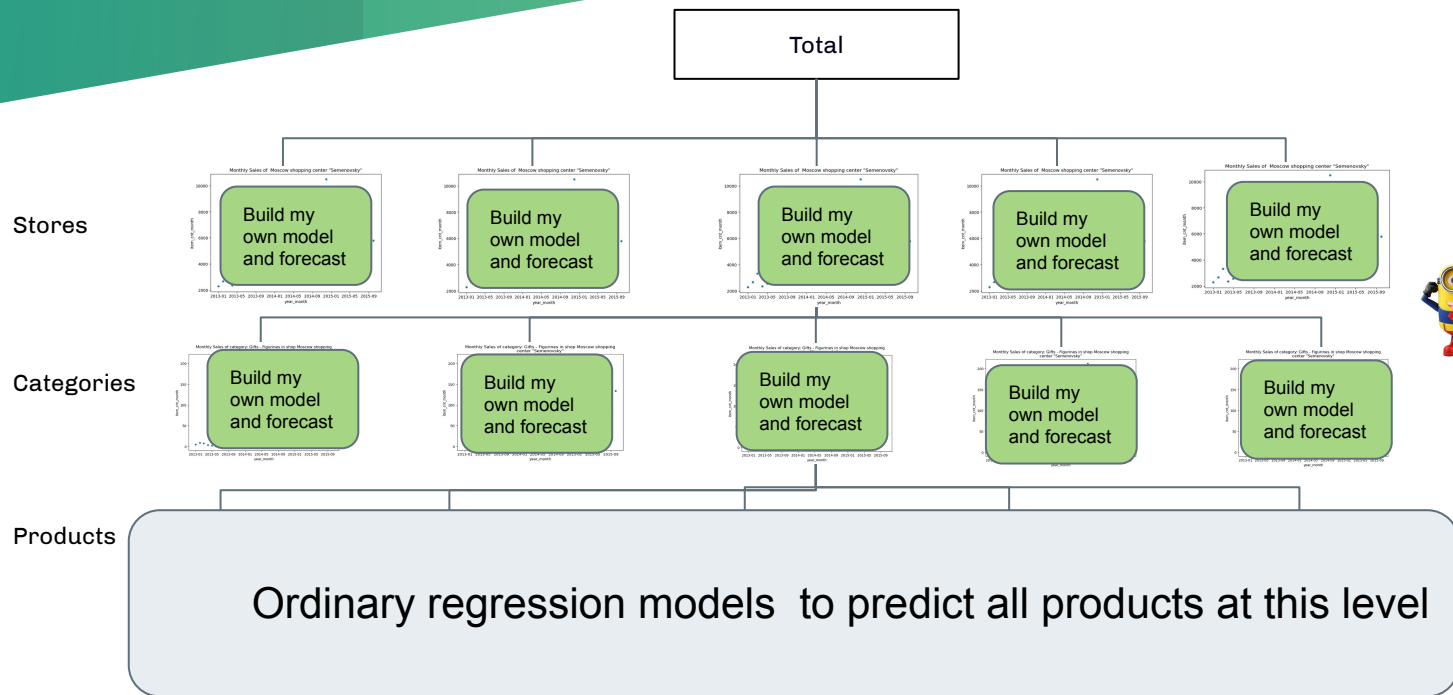


Hierarchical Time Series Forecasting & Reconciliation

- Step 1: Building Hierarchical Time Series:
- **Step 2: Hierarchical Forecasting:**
- Step 3: Forecast reconciliation:



Step 2: Hierarchical Forecasting:



Time Series at category level and store level have noticeable trends and seasonality. We can build time series models on each time series


Without noticeable trends and seasonality, we need to use ordinary regression algorithms such as LightBGM to predict all product sales.

exogenous variable:

- Product name
- Product price
- Lagged data

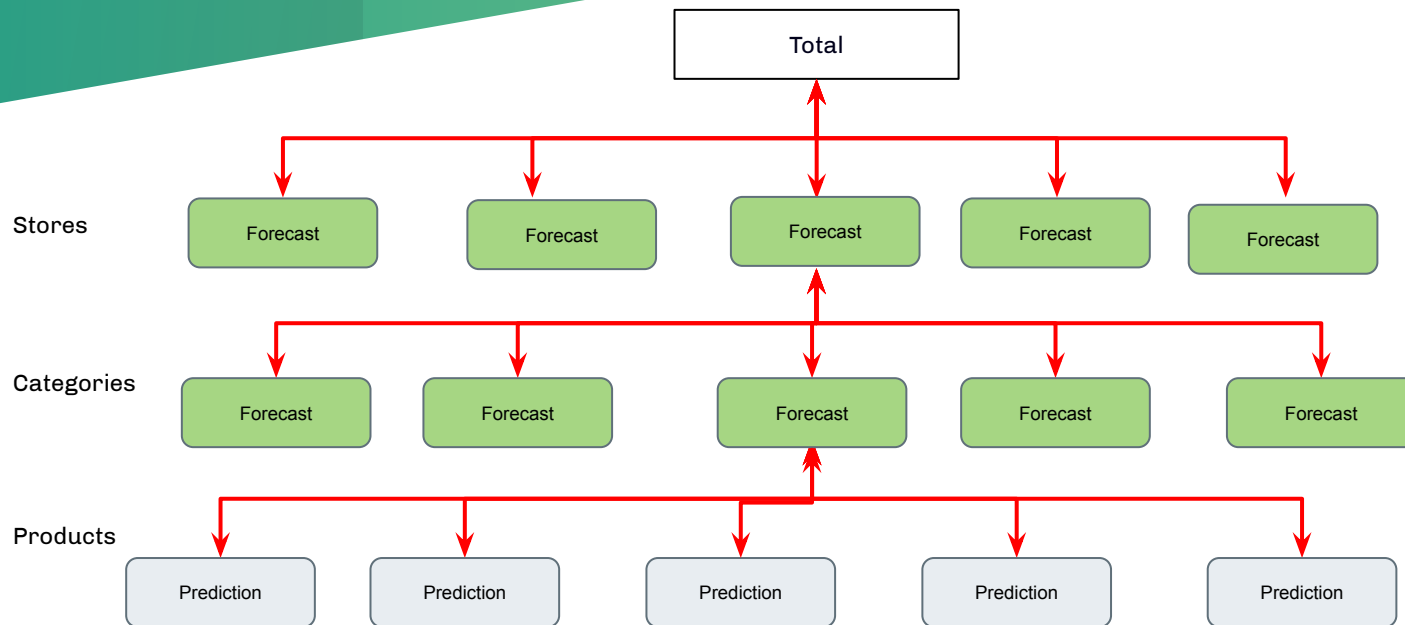
Hierarchical Time Series Forecasting & Reconciliation

- Step 1: Building Hierarchical Time Series:
- Step 2: Hierarchical Forecasting:
- Step 3: Forecast reconciliation:



Step 3: Forecast Reconciliation:

Forecasting Reconciliation



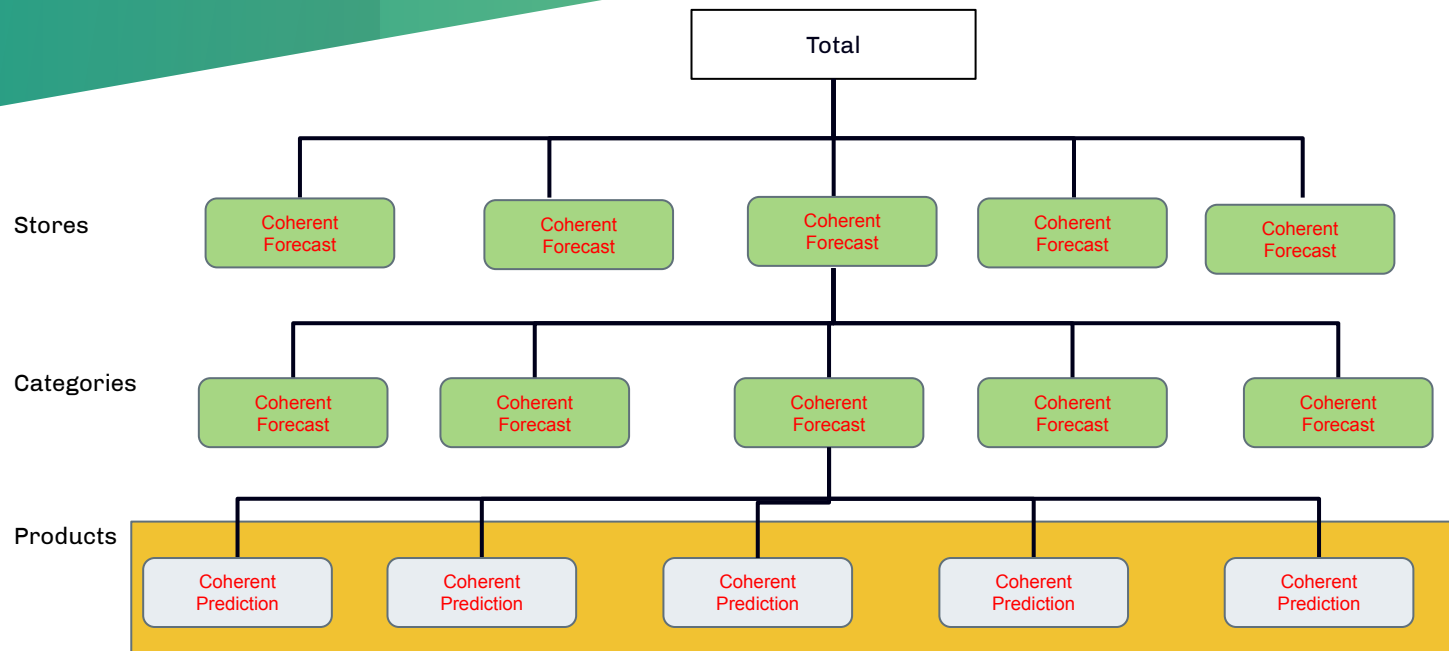
Reconciliation:

- the process of adjusting forecasts to make them coherent and sum up correctly in the hierarchy



Step 3: Forecast Reconciliation:

Done!



Ready to submit

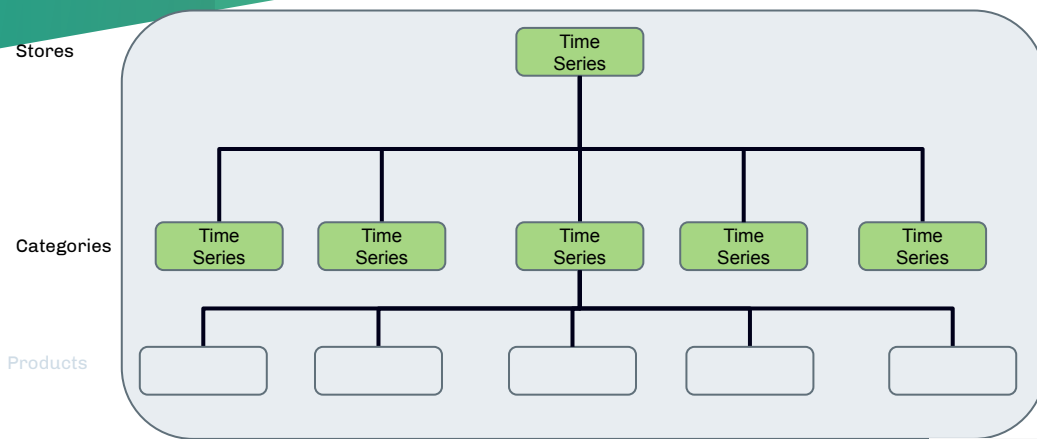


Implementation



- Step 1: Build Hierarchical Time Series
- Step 2a: Top Levels Time Series Modeling and Forecasting
- Step 2b: Bottom level modeling and performance
- Step 3: Forecasting Reconciliation

Implementation of Building Hierarchical Time Series




- Python Package: scikit-hts
- Due to memory constraints, we could only build hierarchical time series up to individual store level
- Sample Code:

We have built 41 such hierarchical trees for 41 stores. There were around 3,444 time series created in total

```
##### create hierachical time series for the shop
#####
level_names = ['item_category_id','item_id']
hierarchy = [['item_category_id']]
df_hier, sum_mat, sum_mat_labels = hts.functions.get_hierarchichal_df(df_oneshop,
                                                                    level_names=level_names,
                                                                    hierarchy=hierarchy,
                                                                    date_colname='year_month',
                                                                    val_colname='item_cnt_month')
```




Implementation of Top Levels Modeling



Time Series at top levels have noticeable trends and seasonality. We can build time series models on each time series

- Python Package: Facebook Prophet
- Took into consideration:
 - yearly-Seasonality
 - Holidays of Russia
- Hyperparameter tuning
 - `Changepoint_prior_scale`
 - `Seasonality_prior_scale`
- Sample Codes:

```
#create model based on best_params
m = Prophet(changepoint_prior_scale=best_params['changepoint_prior_scale'],
            seasonality_prior_scale=best_params['seasonality_prior_scale'],
            yearly_seasonality=True)
m.add_country_holidays(country_name='RU')
m.fit(df)
```

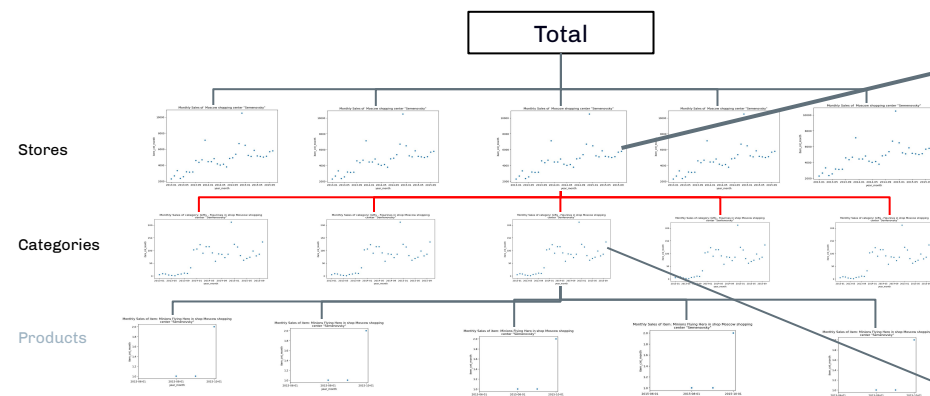
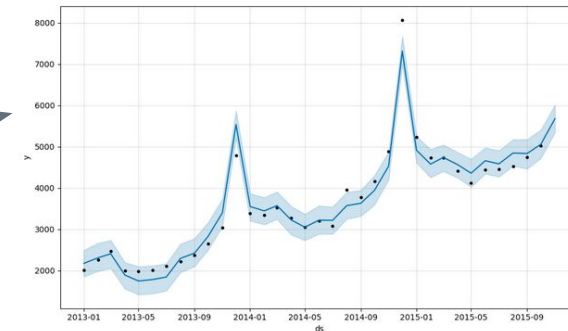


Top levels modeling and performance


forecasting shop: s31

```
getting best preds for forecast_total_sales
{'changepoint_prior_scale': 0.01, 'seasonality_prior_scale': 0.5}
ds      yhat  yhat_lower  yhat_upper
32 2015-09-01 4,835.13    4,465.94    5,178.16
33 2015-10-01 5,057.05    4,718.74    5,407.77
34 2015-11-01 5,676.31    5,345.01    6,008.91
```

Facebook Prophet Forecasting
at store level



```
getting best preds for s31_c72
{'changepoint_prior_scale': 0.2, 'seasonality_prior_scale': 0.01}
ds      yhat  yhat_lower  yhat_upper
32 2015-09-01 101.75    62.76    136.62
33 2015-10-01 103.49    64.70    138.14
34 2015-11-01 103.61    69.35    141.08
```




We have built 3,444 models for
every time series in top levels of
every stores

Facebook Prophet Forecasting
at category level



Implementation of Bottom level modeling and performance




- Intermittent and sparse data:
 - Can not use time series models same as top levels
- Used ordinary regression models:
 - LGBMRegressor
 - XGBRegressor
 - StackingRegressor
- Made use of different exogenous variables to improve prediction:
 - Product name (TfidfVectorizer)
 - Product price
 - Features from category and store
 - Different lag data
 - etc

Bottom level modeling and performance

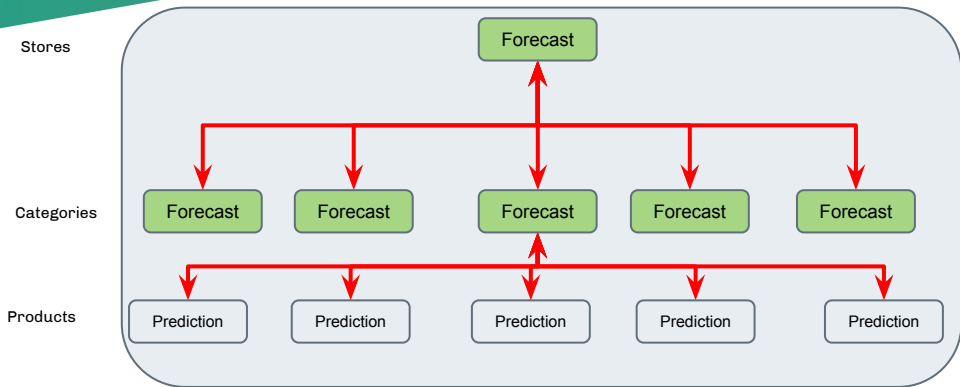
	Training Set (RMSE)	Testing Set (RMSE)	Generalization %	Kaggle Score (RMSE)
LGBMRegressor	1.08	0.93	14.38	0.94407
XGBRegressor	0.97	0.94	3.10	0.96352
StackingRegressor	0.99	0.92	7.17	0.94510

Shap values of model “LGBMRegressor”

- **Shapley value:**
 - the average of the marginal contributions of a feature across all permutations
- **Most important features:**
 - Previous monthly sales of the product in the shop
 - Previous monthly sales of the product in all shops
 - Previous monthly sales of the product category in the shop



Implementation of Forecast Reconciliation



- Python package:
 - scikit-hts
- Reconciliation was carried out on store by store
- reconciliation strategies:
 - Ordinary least squares (OLS) :
 - minimises the total OLS within all coherent forecasts in the hierarchy
 - Structurally weighted least squares (WLSS):
 - minimises the total WLSS within all coherent forecasts in the hierarchy
 - Variance-weighted least squares (WLSV):
 - minimises the total WLSV within all coherent forecasts in the hierarchy

```
##### reconciliation
#####
pred_dict = collections.OrderedDict()
for label in sum_mat_labels:
    pred_dict[label] = pd.DataFrame(data=forecasts[label].values, columns=['yhat'])
revised = hts.functions.optimal_combination(pred_dict, sum_mat, method='WLSS', mse={})
revised_forecasts = pd.DataFrame(data=revised[0:,0:],
                                  index=forecasts.index,
                                  columns=sum_mat_labels)
```

Forecast Reconciliation and Performance

	Before Reconciliation	After Reconciliation					
	Kaggle Score (RMSE)	Reconciliation Strategy: OLS		Reconciliation Strategy: WLSS		Reconciliation Strategy: WLSV	
		Kaggle Score (RMSE)	Improvement %	Kaggle Score (RMSE)	Improvement %	Kaggle Score (RMSE)	Improvement %
LightGBM	0.94407	1.21477	-28.67	0.93976	0.47	-	-
XGBRegressor	0.96352	1.22888	-27.54	0.95546	0.84	-	-
StackingRegressor	0.94510	1.21736	-28.81	0.94063	0.47	-	-



OLS strategy produced many negative results. It made the coherent prediction worse

WLSS strategy produced better results in overall

We encountered some errors when running WLSV strategy

Analysis of Forecast Reconciliation Performance

- The Improvement of reconciliation is not encouraging
- Possible reasons:
 - Due to memory constraints, we just built hierarchical time series up to one store, which may not able to reflect all macro trends or seasonalities
 - Limited hyperparameter tuning
 - Many products/categories have too limited sales transactions
 - Some state-of-the-art reconciliation strategies have not supported by Scikit-hts
 - MinTSample
 - MinTShrink

	Before Reconciliation	After Reconciliation					
	Kaggle Score (RMSE)	Reconciliation Strategy: OLS		Reconciliation Strategy: WLSS		Reconciliation Strategy: WLSV	
		Kaggle Score (RMSE)	Improvement %	Kaggle Score (RMSE)	Improvement %	Kaggle Score (RMSE)	Improvement %
LightGBM	0.94407	1.21477	-28.67	0.93976	0.47	-	-
XGBRegressor	0.96352	1.22888	-27.54	0.95546	0.84	-	-
StackingRegressor	0.94510	1.21736	-28.81	0.94063	0.47	-	-

Conclusions

- Hierarchical Forecasting & Reconciliation did improve the overall prediction of intermittent time series when using correct strategy
- This method can apply to any time series with a hierarchical structure
- Scikit-hts offers a lot of flexibilities
 - Assist to build hierarchical time series from bottom level data
 - We can choose any ML algorithms for different levels of hierarchical time series
 - 3 reconciliation strategies are available at this moment.

Future Improvements

- To build a bigger hierarchy to contain all data
- To try different reconciliation strategies:
 - Variance Scaling
 - MinTSample
 - MinTShrink
- To try different time series models:
 - LSTM
 - N-BEATS
- More feature engineering for bottom level modelling

Q & A

Thank
you!!



Background: 1C company

1C company



A Russian software developer, distributor and publisher based in Moscow.

It is also running a retail chain selling computer software, related services, game consoles, video games, etc.